



Advanced Sample Programs in CompuScope SDKs

Revision 1.04 – 03/18/2021

Copyright © 2021

All Rights Reserved

Gage Contact Information	
Toll Free:	1-800-567-GAGE
Tel:	1-514-633-7447
Fax:	1-514-633-0770
Sales Email:	sales@gage-applied.com
Support Email:	tech-support@gage-applied.com
Web:	http://www.gage-applied.com

Gage is a Product Brand of:



DynamicSignals LLC
900 North State Street
Lockport, Illinois 60441-2200
USA

Toll Free: 1-800-DATA-NOW
Tel: 1-815-838-005
Fax: 1-815-838-4424

<http://www.dynamicsignals.com>

© Copyright DynamicSignals LLC 2006-2021

Tel: 1-800-567-GAGE or +1-514-633-7447. Fax: 1-800-780-8411 or +1-514-633-0770

COMPUSCOPE, GAGESCOPE, AND COMPUGEN are trademarks or registered trademarks of DynamicSignals LLC.

C#, Visual C/C++, .NET, Visual Basic, MS-DOS and Microsoft Windows are trademarks or registered trademarks of Microsoft Corporation. LabVIEW, LabWindows/CVI are trademarks or registered trademarks of National Instruments Corporation. MATLAB is a registered trademark of The MathWorks Inc. Delphi is a trademark or registered trademark of Borland Software Corporation.

Other company and product names mentioned herein may be trademarks or trade names of their respective owners.

Changes are periodically made to the information herein; these changes will be incorporated into new editions of the publication. DynamicSignals LLC may make improvements and/or changes in the products described in this publication at any time.

Copyright © 2006-2021 DynamicSignals LLC. All Rights Reserved, including those to reproduce this publication or parts thereof in any form without permission in writing from Dynamic Signals LLC.

How to reach Gage Product Support

Toll-free phone: (800) 567-4243

Toll-free fax: (800) 780-8411

To reach Gage Product Support from outside North America

Tel: +1-514-633-7447

Fax: +1-514-633-0770

E-mail: prodinfo@gage-applied.com

Web site: <http://www.gage-applied.com>

On-line Support Request Form: <http://www.gage-applied.com/support/support-form.php>

Table of Contents

OVERVIEW 4

OPTIONAL FIRMWARE IMAGES 5

STREAMING SAMPLE PROGRAMS 5

GAGESIGNALAVERAGING 10

GAGECSPRF 11

GAGEFASTACQUIRE2DISK 12

GAGEEVENTS, GAGECALLBACK..... 16

TECHNICAL SUPPORT 17

Overview

There are three Gage Software Development Kits (SDKs) for user programming of Gage CompuScope cards: one for C/C#, MATLAB, and LabVIEW. Each of these SDKs includes advanced sample programs that are provided in a separate sub-folder and that are not described within the standard CompuScope SDK documentation. These Advanced sample programs are described within this document. This document describes the advanced sample programs in a generic fashion that is not specific to any of the three SDKs.

The table below lists which sample programs are supported under which programming language. Some sample programs for VisualBasic.NET and Delphi are provided within the C/C# SDK. The user may use the C Sample programs as a guide to construct sample programs that are unavailable for VB.NET and Delphi. Usage of some advanced sample programs requires installation of optional eXpert firmware. These programs are indicated by an asterisk (*).

To determine which CompuScope models are supported by which advanced sample programs, please check the online Advanced Functionality Matrix at:

http://www.gage-applied.com/products/eXpert_FPGA_technology/eXpert_compatibility_matrix.htm

	C	C#	VB .Net	LabVIEW	MATLAB	CVI	Delphi
GageAcquire	X	X	X	X	X	X	X
GageComplexTrigger	X	X		X	X	X	
GageDeepAcquisition	X	X		X	X	X	
GageMultipleRecords	X	X		X	X	X	
GageMultipleSystems	X	X		X	X	X	
Streaming Sample Programs* (GageStream2Analysis, GageStream2Disk ...)	X						
GageSignalAveraging*	X	X		X	X	X	
GageCsPrf	X	X				X	
GageFastAcquire2Disk	X	X		X	X		
GageCallback	X					X	
GageEvents	X	X				X	

Optional Firmware images

Some CompuScope models have the ability to be reconfigured with optional alternative firmware allowing on-board processing of waveform data before they are transferred to PC RAM. Currently, these firmware options include eXpert Multiple Record Signal Averaging, and eXpert Data Streaming (PCI Express CompuScopes only). When a CompuScope is updated with an optional firmware image, the image information is stored in CompuScope non-volatile memory. This memory has space for two or three firmware images: the standard CompuScope operating image and up to two optional images. The contents of non-volatile memory may be queried by an application using the `CsGet(hSystem, CS_PARAMS, CS_EXTENDED_OPTIONS, &i64ExOptions)` call, where `hSystem` is the CompuScope system handle. `CS_PARAMS` and `CS_EXTENDED_OPTIONS` are constants defined in `CsDefines.h` and `CsExpert.h` respectively. `i64ExOption` is a 64-bit integer type variable that is filled with the results of the query. The lower 32 bits of the `i64ExOption` will contain information about first alternative image and the higher 32 bits will contain information about the second one. There are corresponding calls to query available firmware in MATLAB (`CsMI_GetExtendedOptions.m`) and LabVIEW (`CsLv_GetExtendedOptions.vi`).

Based on the information about available firmware images, the user application can decide which image to load. Firmware images are loaded from any SDK by bitwise ORing the CompuScope mode (1 for “Single”, 2 for “Dual”, 4 for “Quad” and 8 for “Octal”) with a specific constant that indicates the image number. For instance, to specify an alternative image from C, either the constants `CS_MODE_USER1` or `CS_MODE_USER2` (for images #1 or #2) should be bitwise ORed with the `u32Mode` member of `CSACQUISITIONCONFIG` structure before it is used in the `CsSet()` call. (Note that the firmware image is not actually loaded until a call is made to `CsCommit()`).

Each of the CompuScope SDKs (C/C#, MATLAB and LabVIEW) provides a programming example for each optional firmware image (currently FIR filtering and signal averaging). The programming sequence for the loading each image, which is described above for C, is illustrated within each programming example

Streaming Sample Programs

Introduction

Gage PCI Express CompuScopes now support direct streaming of waveform data to PC RAM for display analysis or storage. Streaming is an alternative to conventional Memory Mode operation. In Memory Mode, waveform data are first acquired into on-board CompuScope memory and are later downloaded to PC RAM in a separate operation. Accordingly, at any given time in Memory Mode, on-board memory is either written-to or read-from – never both. By contrast, Gage PCI Express CompuScopes are equipped with dual-port memory that allows on-board memory to be simultaneously written-to and read-from. Using this architecture, a data stream may be established that flows continuously through the on-board memory, which is used as a large FIFO (*First In First Out*).

The CompuScope C/C# SDK includes several sample programs that illustrate data streaming using PCI Express CompuScopes. In order for these sample programs to operate, the PCI Express CompuScope hardware must be equipped with the *eXpert Data Streaming* option. If your PCI Express CompuScope is not equipped with this eXpert Data Streaming, please contact your Gage Sales Agent about procuring this option, which may be deployed remotely with no factory upgrade required.

This document provides an overview and description of the operation of the streaming sample programs, as well as a description of the four new CompuScope API methods that are used within these sample programs to implement streaming.

CompuScope Streaming Operation

The CompuScope C/C# SDK includes four separate sample programs that illustrate streaming usage. These sample programs are listed in the table below:

Sample Program	Description	Comments
<i>GageStream2Analysis_Simple</i>	Simple, single-threaded streaming acquisition from a single CompuScope	Best starting point for single CompuScope streaming software application. A simple example analysis is performed upon a raw waveform data stream.
<i>GageStream2Analysis</i>	Multi- threaded streaming acquisition from a single CompuScope or Master/Slave CompuScope System	Streams waveform data, which is ultimately not stored. Analysis is performed on time-stamp data and resultant values are stored to hard drive.
<i>GageStream2Disk</i>	Multi- threaded streaming acquisition and storage from a single CompuScope or Master/Slave CompuScope System	Streams waveform data to hard drive. Use this sample for applications that must stream and save a huge amount of data. Sustaining this stream requires a PC with large and fast hard disk drives.
<i>GageStream2RAM</i>	Multi-threaded streaming acquisition and storage from a single CompuScope	Streams waveform data to a large PC RAM memory buffer and stores these data afterwards. Use this sample for a PC that do not have fast hard disk drives and applications that need to capture and save an amount of data that is bigger than the CompuScope on-board memory but smaller than the PC RAM.

The first two sample programs are designed to serve as starting points for user software applications and include the word “analysis”, which is used in the most general sense and could refer to any software process that consumes the data stream. This process could be actual numerical data analysis or other processes, such as data display or data transfer to a separate device. The analysis algorithm is clearly separated from the main body of the source code and so may easily be replaced with your own data analysis algorithm.

The simplest streaming sample program, *GageStream2Analysis_Simple*, operates only a single CompuScope card. It is simple because all programming is done within a single software thread. The example analysis algorithm within *GageStream2Analysis_Simple* simply sums up the contents of the waveform data buffer for all channels and stores the resultant summed value to an ASCII data. While, by itself, it is not useful in any real-world application, this analysis algorithm was chosen for its simplicity.

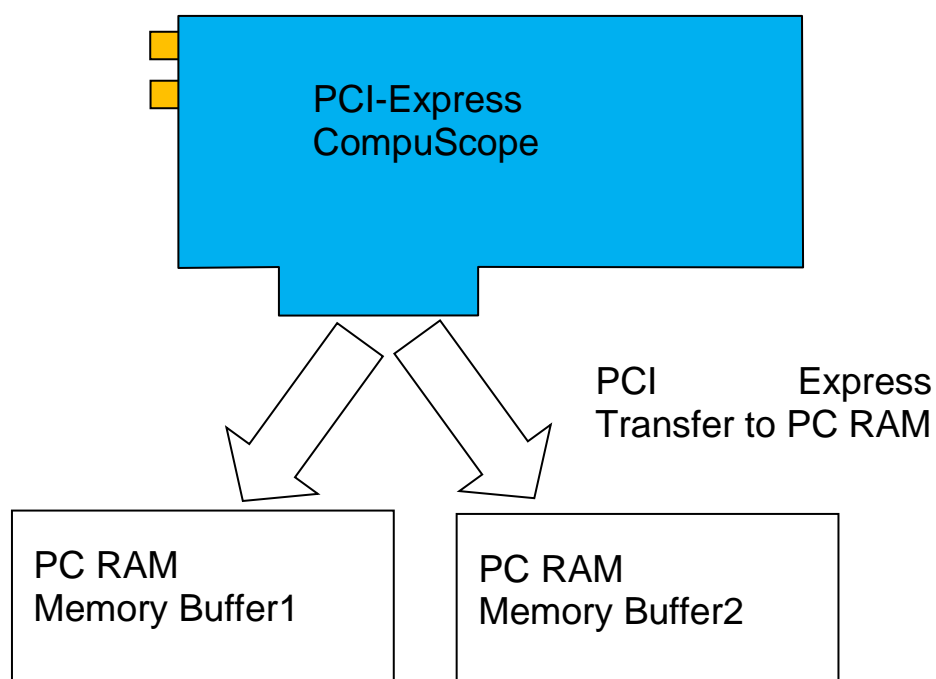
The more complex sample program, *GageStream2Analysis* operates a single CompuScope card or a Master/Slave CompuScope system. The data stream from each CompuScope card is controlled from a separate software thread. Unlike *GageStream2Analysis_Simple*, which analyzes the raw waveform data, *GageStream2Analysis* analyzes the time-stamp values associated with each waveform. *GageStream2Analysis* is the best starting point for a streaming software application that operates multiple CompuScopes.

The third sample program, *GageStream2Disk* streams data from a single CompuScope card or from a Master/Slave CompuScope system to a target file on the hard drive. The data stream from each CompuScope card is controlled from a separate software thread and is saved into the hard drive while the streaming acquisition is in progress. Data are stored in a header-less binary file, which consists of the waveform data in binary followed by the “tail” for each record, which contains the associated time-stamp value.

Similar to *GageStream2Disk*, *GageStream2RAM* streams also data from a single CompuScope to a target file on the hard drive. However, while acquisition is in progress, data received from CompuScope card is copied to a huge RAM buffer allocated previously in application. All data is saved into the hard drive only at the end of acquisition.

The general method of streaming waveform data from a PCI Express CompuScope to PC RAM is illustrated in the diagram below.

Using the ***CsStmAllocateBuffer()*** method, the user first allocates two PC RAM contiguous buffers of identical size, called *Buffer1* and *Buffer2*, that serve as targets for the data stream. If the operating system is unable to allocate the requested contiguous buffer, an allocation error will occur. The user toggles the two buffer targets as each one gets successively filled. While data are streaming to one buffer, the user analyzes (consumes) waveform data from the second buffer. As long as the user is able to sustainably analyze (or consume) all waveform data in the buffer under analysis before the streaming target buffer is filled up with new waveform data, then the streaming acquisition can proceed indefinitely with no data loss.



The size of the target memory buffer is internally manipulated so that it is simply related to the *SegmentSize*. Specifically, if the requested buffer size is less than the *SegmentSize*, then the buffer size is internally increased until the *SegmentSize* is an integral multiple of it. This way, the full segment may be transferred to the buffer in an integral number of data transfers and no partial transfers are required. On the other hand, if the requested buffer size is greater than the *SegmentSize*, then the buffer size is increased until it is an integral multiple of *SegmentSize*. This way, the buffer will fit an integral number of segments and no partial segments occur. These internal buffer size adjustments were made in order to keep the overall scheme simple.

Optimal selection of the buffer size depends upon the user application. Let us consider two examples of the acquisition of multiple short segments. In lightning monitoring, where lightning triggers occur randomly, the user would not want the buffer size to exceed *SegmentSize*. Otherwise, the consuming analysis process might have to wait hours or days before a lightning trigger occurred to fill up the buffer and allow consumption to begin. In lightning test, therefore, the user should make the buffer size equal to the *SegmentSize* so that lightning waveforms may be consumed as quickly as possible. Now consider the case of a scanning application where the triggers occur at a regular fast frequency. In this case, the user knows that multiple triggers will come and they can afford to make the buffers size much larger than the *SegmentSize*. If they had made the two sizes equal, then the software would waste a lot of time toggling between the two buffers, which would compromise overall throughput. If the user wanted to do only post-processing, then they should make the buffer size as large as possible in order to maximize throughput. If, on the other hand, they wanted to display some results during the scan, then they would have to choose an intermediate buffer size that allowed for decent throughput while still providing a reasonable result update rate.

Within the streaming sample programs after initial configuration, acquisition on CompuScope hardware is initiated by a call to ***CsDo(hSystem, ACTION_START)***, as in normal Memory Mode. This call begins acquisition of waveform data into on-board CompuScope memory, which acts like a large FIFO. Next, the user initiates transfer of data from on-board CompuScope memory to one of the streaming target buffers using the ***CsStmTransferToBuffer()*** method. The waveform data within the other filled data buffer are then analyzed (using the *SumBufferData()* method, for example, for *GageStream2Analysis_Simple*). After this analysis is complete, the user then awaits for the completion of data transfer to the first buffer using the ***CsStmGetTransferStatus()***. The status returned by this method is checked to see if there was a FIFO overflow, which will terminate the sample program. If no overflow has occurred, then the sample program swaps the two buffers and continues the process.

Upon detection of any FIFO overflow errors (and the associated data loss), the Gage streaming sample programs will terminate after displaying the FIFO overflow error message. This need not be the case and the user can program his streaming application to continue acquiring after detection of a FIFO overflow error. We strongly recommend, however, that the user's software application keep track of any occurrences of the FIFO error and, therefore, any potential data loss. Please note that data loss may only occur after the two buffers have been filled at least once.

While the Gage sample programs use two toggling buffers, the user could modify the program to use more. Using a number of buffers that is equal to the number of cores on a multi-core CPU may allow a multi-threaded application to simultaneously analyze multiple buffers in a fashion that would improve the overall analysis throughput.

Used within the *GageStream2Analysis_Simple* sample program, the *SumBufferData()* method is an example analysis routine that sums all points within the data buffer. The resulting sum values are stored in an ASCII data file by *GageStream2Analysis_Simple*. The summing analysis is not useful by itself but was chosen for its simplicity and fast execution speed. If the user chooses *GageStream2Analysis_Simple* as a starting point for their own software application, the user should replace the contents of *SumBufferData()* with his own analysis

subroutine. Again, this analysis routine could be data display, transfer, storage or any operation that consumes the data stream. Naturally, more complex analysis routines may require more processing time. Sufficiently complex data analysis will retard the analysis process to the point where the buffer under analysis cannot be emptied before the stream target buffer is filled, which will lead to a FIFO overflow and data loss. The complexity of the analysis that the streaming process can handle will depend upon system configuration (CPU core count and speed, PC RAM volume, etc.). For best performance, no other software applications should be executed while using eXpert data streaming.

The streaming sample programs operate in two distinct modes *Continuous Mode* and *Segmented Mode*. In Continuous Mode, the user does one long continuous acquisition that is initiated by a single trigger. For Continuous Mode operation, the user must specify a *SegmentCount* of 1 within the INI file. The *Depth* and *SegmentSize* values should correspond to the number of samples that the user wants to acquire for the one large waveform, as in the *GageAcquire* project.

In Segmented Mode, which is the streaming version of Multiple Record Mode, the user specifies a *SegmentCount* greater than 1 and must ensure that a corresponding number triggers occur. Otherwise, the program will appear to freeze as it awaits trigger events that never occur (assuming that the Trigger Timeout is infinite). Each waveform is transferred to PC RAM with a *Tail* that contains Trigger TimeStamp information.

As with other C/C# SDK sample programs, configuration parameters for the streaming sample programs are entered in an INI text file. Both *GageStream2Analysis_Simple* and *GageStream2Analysis* obtain their configurations from the same INI text file, while *GageStream2Disk* has its own INI file. Most parameters have the same meaning as in other CompuScope C/C# SDK sample projects.

There are maximum size limits on the *SegmentSize* and *SegmentCount* for a streaming acquisition. The maximum value of *SegmentSize* depends on CompuScope model and it can be retrieved using the API `CsGetSystemCaps()` with the parameter `CAPS_STM_MAX_SEGMENTSIZE`. The returned value from the API must be divided by the number of active channels and the result is the maximum segment size that can be used for data streaming in multiple record mode. Beyond this limit, only single record, infinite data streaming is supported. To achieve this, the user sets the *SegmentSize* to -1, the acquisition will be endless and the user must terminate it by aborting the acquisition.

The maximum allowed 32-bit *SegmentCount* value is 4 Billion (actually 4,294,967,295) records. It is possible to use higher *SegmentCount* values because the API allows for construction of a 64-bit *SegmentCount* variable. The user would have to write their own software to exploit this, however.

By necessity, waveform data transferred during a streaming acquisition must include data from all active channels. The number of active channels is determined by the CompuScope channel mode (Octal, Quad, Dual or Single). There is currently no provision for streaming only data from a subset of the active channels.

The list below illustrates how data from different channels are organized in memory depending on the channel mode (Octal, Quad, Dual or Single). Each sample in memory is represented as $S_n i$, where n is channel number, and the i is the sample index. Example: $S_3 0$ represents the first sample of the channel 3 and $S_5 3$ the fourth sample of the channel 5.

Octal (8 channels)

$S_{10} S_{20} S_{30} S_{40} S_{50} S_{60} S_{70} S_{80} S_{11} S_{21} S_{31} S_{41} S_{51} S_{61} S_{71} S_{81} \dots$

Quad (4 channels)

$S_{10} S_{20} S_{30} S_{40} S_{11} S_{21} S_{31} S_{41} S_{12} S_{22} S_{32} S_{42} \dots$

Dual (2 channels)

$S_{10} S_{20} S_{11} S_{21} S_{12} S_{22} S_{13} S_{23} S_{14} S_{24} \dots$

Single (1 channels)

$S_{10} S_{12} S_{13} S_{14} S_{15} S_{16} S_{17} S_{18} S_{19} S_{110} \dots$

GageSignalAveraging

Usage of the signal averaging optional firmware image allows repetitive waveform acquisitions to be rapidly averaged, in order to reduce random noise. In the past, signal averaging required waveforms to be downloaded for averaging within the host PC's CPU so that averaging was limited by the data transfer speed. With the SignalAveraging firmware, repetitive waveforms are averaged within the firmware with no data transfer required until up to 1024 averages have been performed. Consequently, much higher repetition rates may be achieved.

Once the signal averaging firmware has been loaded, the user adjusts the Number of Averages to be acquired for each waveform to be averaged. In addition, the user may set the SegmentCount to be greater than 1 so that averaged waveforms are stacked in on-board acquisition memory, as in Multiple Record mode. In fact, the firmware does not actually average waveforms but *co-adds* them, which means that the waveforms are summed together but the resultant summed waveform is not divided by the number of averaged waveforms in order to calculate the true signal average. The firmware performs co-adding rather than true averaging because division is a complex operation within an FPGA, is often unnecessary and may easily be done in software, if necessary. With the SignalAveraging firmware loaded, the CompuScope hardware co-adds a preset number of consecutive waveforms, and then stacks the resultant co-added waveforms in on-board memory, as is usually done in Multiple Record Mode.

As an example, consider the acquisition of 10,000 sample waveforms using SignalAveraging in which the Number of Averages was set to 1000 and the SegmentCount was set to 2000. In this example, a total of $2000 \times 1000 = 2$ million waveforms would be acquired by the CompuScope hardware. However, each consecutive 1000 waveforms would be co-added together and the resultant 2000 co-added waveforms would be stacked in on-board memory. Since the hardware re-arm between successive groups of 1000 averaged acquisition is very fast, the user could use this functionality to follow noisy signals that rapidly evolve in time. Alternatively, the user could co-add the 2000 waveforms in software in order to create a "Super-averaged" waveform that is the result of 2 million averages.

Co-added waveform data are stored within a 32-bit format buffer within the on-board memory. The resulting averaged waveform is therefore transferred as a 32-bit data buffer. With the SignalAveraging firmware loaded, the Sample Size, Sample Resolution and Sample Offset values are changed to reflect the 32-bit data format. Querying these parameters after the firmware is loaded will return the updated values. The updated Sample Resolution includes the number of Number of Averages factor by which the co-added waveforms must be

divided in order to obtain the averaged waveform. Consequently, the updated Sample Resolution and Sample Offset values may be used with the Voltage Conversion equation in order to directly convert the co-added waveforms into for the correct voltage waveforms.

Each SDK contains an advanced sample program that uploads the SignalAveraging image, performs multiple averaging acquisitions and then displays or stores the resulting averaged waveforms.

GageCsPrf

GageCsPrf is an advanced sample program that may be used to evaluate the repetitive capture performance of CompuScope hardware. The program makes a series of repetitive acquisitions using the specified settings and provides subsequent timing measurements. All timing measurements are done using the QueryPerformanceCounter() Windows API timing function. The total time required to complete a single acquisition is provided along with its inverse, the Pulse Repeat Frequency (PRF). In addition, the timings of the separate operations that occur within one acquisition are provided. Repetitive acquisition sequences are repeated using different acquisition depths and timing results are provided for acquisition depths.

GageCsPrf is based upon the GageAcquire sample program and uses similar controls. Only controls that are different from those of GageAcquire are described here. See the GageAcquire documentation for the common controls.

Since GageCsPrf does not store acquired data, the *SaveFileName* key within the INI file is ignored. The timing results file name is specified by the *ResultsFile* key in the *PrfConfig* branch in the INI file. Also, the *TransferLength* key in the *Application* branch, along with the *SegmentSize* and *Depth* keys in the *Acquisition* branch, are ignored since the *Depth* is internally adjusted by *GageCsPrf*. Internally, the *Depth* is always selected as a power of two. The range of internally selected *Depth* values is bound by the smallest power of two that is greater than or equal to the value of the *StartDepth* key in the *PrfConfig* branch and the highest power of two that does not exceed the value of the *FinishDepth* key in the *PrfConfig* branch.

Results are stored in a tab delimited text file. The first 5 lines describe the measurement configuration, such as the CompuScope model number and memory size, acquisition configuration and number of acquisitions in one repetitive capture sequence (loop count). These lines are followed by 7 columns of data. The first column (*depth*) specifies the size of the acquisition. The remaining columns are the result of the timing measurements. *Total time* is the time required by the complete acquisition sequence and *PRF* is its inverse, the repeat frequency of acquisitions. *Start time* is the time required for execution of the CsDo (ACTION_START) method. In this time the CompuScope system will perform all necessary operations to start an acquisition. *Busy time* is the time taken by the data acquisition itself. Please note that this time includes waiting for the trigger event. Consequently, if the trigger event is infrequent, the measured *Busy time* will be long. For measurement of the maximum PRF, trigger from a fast signal source, such as a sine wave with a frequency of 1 MHz or more. *Transfer time* and *transfer rate* describe data transfer from one channel. Please note that the data *transfer time* includes two components: a fixed transfer set-up overhead time and the actual data transfer duration, which is proportional to the data volume. As data volume increases, the importance of the overhead time diminishes. Consequently, the calculated aggregated *transfer rate* improves with the *Depth*.

Within GageCsPrf, the power-saving mode is enabled. This is fine for all CompuScope models except the CS82G and CS8500. For these models, power-saving mode will severely reduce repetitive capture performance (PRF). In order to get the best PRF from these models, you must disable power-saving mode. This is done by adding 128 to the Mode value within the PRF.ini file. For single-channel mode, use "Mode=129" and for dual-channel mode, use "Mode=130".

By following the coding illustrated in GageCsPrf, a user can achieve the fastest possible repetitive capture performance from CompuScope hardware.

GageFastAcquire2Disk

Several Gage customers require their CompuScopes to acquire fast repetitive waveforms and store them as quickly as possible to hard drive. CompuScope hardware allows for fast PCI data transfer, extremely fast hard drive arrays are commercially available and MS Windows provides powerful multi-threading capability. Optimally combining these elements, however, to fully exploit their performance is not straightforward. In order to allow Gage customers to realize the best possible performance, Gage has created *GageFastAcquire2Disk*. This sample program uses a flexible, CompuScope driver-level engine and that allows the fastest possible repetitive CompuScope capture to hard drive.

A GageFastAcquire2Disk sample is now available as an Advanced Sample Program within CompuScope SDKs for C/C#, LabVIEW and MATLAB. The program will support waveform acquisition and storage sequence from any single CompuScope system (single card or Master/Slave system) constructed using any CompuScope model. User may even launch concurrent instances of Gage FastAcquire2Disk in order to simultaneously acquire data from independent CompuScope systems. The sample programs are all relatively thin shells that call the GageFastAcquire2Disk engine, which operates from within the CompuScope drivers. Performance is therefore virtually identical from all three SDKs. The engine incorporates the following features to maximize overall throughput:

- Multi-threading for concurrent, parallel operations where possible – for example, storing previously acquired waveforms to hard drive while new waveforms are being acquired
- Pre-opening of target waveform files to avoid costly file open operations during acquisition sequence
- Block transfer size from CompuScope hardware optimized for most efficient usage of PC RAM under Windows
- Optimized storage file transfer size for fast storage rate and easy management under Windows
- Hardware interrupt usage to avoid costly repetitive polling operations

The GageFastAcquire2Disk engine uses a lot of system resources. While CPU processing power per se is not important, the host system should be equipped a fast PCI bus, a fast memory bus, a fast hard drive and as much PC RAM as possible. Specifically, the PC systems should be equipped at least an amount of PC RAM that is equal to the number of active CompuScope channels multiplied by the 128 MB. Further, before running any GageFastAcquire2Disk sample program, the user should close all unnecessary Windows applications, since these will divert resources from GageFastAcquire2Disk.

From the GageFastAcquire2Disk sample program all SDKs, the Gage FastAcquire2Disk operation method is the same. Before activating the GageFastAcquire2Disk engine, the sample program must first initialize the drivers and commit CompuScope configuration settings, as usual. Next, the program activates the GageFastAcquire2Disk engine and initializes it with acquisition settings (e.g. number of waveforms to be acquired, data of interest that is to be transferred). The initialization also validates input parameter and allocates buffers required for GageFastAcquire2Disk acquisitions. Finally, initialization creates all Windows folders and pre-opens all SIG files that will be required for the acquisitions. Next, the program starts the GageFastAcquire2Disk acquisition sequence, so that the CompuScope hardware begins acquiring waveforms and

storing them to files. Normally, the process continues until the requested number of waveforms has been acquired and stored to hard drive. From C, all control of the GageFastAcquire2Disk engine is invoked using the CsExpertCall() API call with different u32ActionId constants. The user may issue the following commands and queries:

1. COMMAND INITIALIZE (implemented by the CsExpertCall() with u32ActionId = EXFN_DISK_STREAM_INITIALIZE): Activate and initialize the GageFastAcquire2Disk engine. Allocate required resources and validates input GageFastAcquire2Disk parameters. Create target folders and pre-open target SIG files.
2. COMMAND START (implemented by the CsExpertCall() with u32ActionId = EXFN_DISK_STREAM_START): Start the GageFastAcquire2Disk acquisition sequence, which will normally terminate only once all waveforms are acquired and stored to hard drive files.
3. COMMAND STOP (implemented by the CsExpertCall() with u32ActionId = EXFN_DISK_STREAM_STOP): Stop the GageFastAcquire2Disk acquisition sequence after completing any acquisition that is in progress and storing all resultant waveform files to hard drive.
4. COMMAND CLOSE (implemented by the CsExpertCall() with u32ActionId = EXFN_DISK_STREAM_CLOSE): De-activate the GageFastAcquire2Disk engine. Free all allocated system resources, including active CompuScope handles.
5. QUERY STATUS (implemented by the CsExpertCall() with u32ActionId = EXFN_DISK_STREAM_STATUS): Returns status that indicates if GageFastAcquire2Disk acquisition sequence is finished.
6. QUERY # ACQUISITIONS (implemented by the CsExpertCall() with u32ActionId = EXFN_DISK_STREAM_ACQ_COUNT): Return the number of acquisitions that have been completed since the GageFastAcquire2Disk acquisition sequence has started.
7. QUERY # FILES (implemented by the CsExpertCall() with u32ActionId = EXFN_DISK_STREAM_WRITE_COUNT): Return the number of files that have been stored since the GageFastAcquire2Disk acquisition sequence has started.

As discussed, adjustment of configuration settings (e.g. sampling rate, input range) for GageFastAcquire2Disk is done in a similar fashion as for the standard SDK sample programs. However, there are some additional setup variables that have been added for GageFastAcquire2Disk. These variables are described below only for FastAcquire2Disk.INI configuration file for the C Sample FastAcquire2Disk program but corresponding variables exist for the other SDKs.

Within FastAcquire2Disk.INI, the *AcqCount* variable determines the total number of acquisitions that will be performed by GageFastAcquire2Disk. Only completion of this number of acquisition will terminate the GageFastAcquire2Disk engine, unless it encounters an error or is intentionally stopped or aborted. The *Channels* variable selects the channels whose data will be stored to hard drive files. For instance, if a user is only interested in the even channel on an 8-channel CompuScope card, he can set *Channels=2,4,6,8*.

As in the standard SDK samples, GageFastAcquire2Disk allows adjustment of the *StartPosition* and *TransferLength*, which determine how much waveform data are to be downloaded from which point in the waveforms. In addition, GageFastAcquire2Disk allows the user to select how many records to download and from which starting record for Multiple Record acquisitions. For instance, if the user does Multiple Record acquisitions acquiring 120 Records, he can choose to download and store only 10 records starting from record #40 by setting *RecordStart=40* and *RecordCount=10*.

Finally, there are three separate Trigger TimeOut Settings in GageFastAcquire2Disk, all of which are expressed in milliseconds and all of which are disabled by a value of -1. The *TriggerTimeOut* in the *Acquisition* key within the INI file, which is standard for all C SDK programs, is a hardware Trigger TimeOut. There is a high-speed counter on the CompuScope hardware that forces a trigger event if the counter value corresponding to the Trigger TimeOut is exceeded. The advantage of this hardware Trigger TimeOut is that it has sub-microsecond precision and is also very accurate and deterministic, as compared to a software Trigger TimeOut. The disadvantage, however, is that the hardware counter will eventually roll over, since it has finite width.

The *TriggerTimeOut* in the *AppConfig* key is a software Trigger TimeOut. Its use is necessary when the required Trigger TimeOut exceeds the rollover limit of the hardware Trigger TimeOut (typically over an hour to several hours, depending upon the CompuScope model and sampling rate). When using this software Trigger TimeOut, the user should disable the hardware Trigger TimeOut by setting it to -1.

The final timeout, *StatusTimeout*, is an intentional software latency introduced so that the QUERY operation listed above (QUERY STATUS, QUERY # ACQUISITIONS and QUERY # FILES) are not called in a loop at a rate sufficient to compromise the speed of the GageFastAcquire2Disk data transfer and storage operations. The default setting of 500 milliseconds ensures that the driver will not poll the status of GageFastAcquire2Disk more than twice a second, which will not significantly affect operation and is usually a sufficient status update rate for the user.

Gage FastAcquire2Disk stores waveforms acquired by CompuScope hardware in GageScope SIG files, which are binary data files with a 512 byte information header. For easy management under Windows, Gage has limited the maximum SIG files size created by GageFastAcquire2Disk to a maximum of 256 Megabytes. (This number may be increased in future, but the principles discussed below will still apply). In order to optimize system memory usage, GageFastAcquire2Disk internally limits the size of a single block data transfer from the CompuScope hardware to a maximum of 128 Megabytes per active channel. Consequently, the user should have a bare minimum of 128 Megabytes of PC RAM per active channel. Of course, this bare minimum does not leave any room for the O/S or for applications, so that significantly more than this minimum is highly recommended for optional operation.

GageFastAcquire2Disk stores all acquired waveform data within GageScope SIG files that reside within a folder structure designed for easy management. The folder structure is created upon initialization of the GageFastAcquire2Disk engine. First, all files are placed within a Master folder local to the GageFastAcquire2Disk application, whose name is specified by the *FolderName* variable (Default= "Signal Files"). Next, folders within the Master folder are created for data from different active CompuScope channels. These folders are given names in the form YYYY-MM-DD_HH-MM-SS_CHANXX, where YYYY-MM-DD and HH-MM-SS are the system date and time when the folders are created upon initialization and XX is the CompuScope channel number. Within these folders, a series of sequential folders called *Folder.001*, *Folder.002*, *Folder.003*, etc. are created. The contents of each of these folders is limited to a maximum of 16,000 SIG files and initialization creates a sufficient number of folders to accommodate the entire GageFastAcquire2Disk acquisition sequence. Within these folders, SIG files called *File-00000.SIG*, *File-00001.SIG*, etc. are created and opened so that they are ready to accept data from the subsequent GageFastAcquire2Disk acquisition sequence. We describe below the three most common GageFastAcquire2Disk use cases and describe how the resultant waveform data are stored in each case.

The most common use case is acquisition of many relatively short waveforms (a few hundred Samples ~ 100 MegaSamples) by CompuScope hardware in Single Record Mode. In this case, the GageFastAcquire2Disk engine will generally pack many single waveforms into a single Gage SIG file. All but the last SIG file will contain the same number of records. The last SIG file will most likely contain a smaller number of left-over records. As an example, consider the user who configures GageFastAcquire2Disk to acquire twenty thousand waveforms of

65536 Samples each in Single Record Mode from a 14-bit CompuScope, where each Sample will occupy 2 data bytes. In this case, GageFastAcquire2Disk will fill each 256 MB SIG File with 2048 records, since $(256 \times 1024 \times 1024 \text{ MB}) / (65536 \text{ Samples/waveform}) / (2 \text{ Bytes/Sample}) = 2048 \text{ waveforms}$. GageFastAcquire2Disk would create nine such SIG files and the final tenth SIG file would contain only the remaining $20,000 - 9 \times 2048 = 1568$ records.

The second case is acquisition of relatively long waveforms (>256 Megabytes) by CompuScope hardware in Single Record Mode. Since the waveform size exceeds the maximum file size, the GageFastAcquire2Disk engine will divide the waveform into multiple contiguous SIG files. For example, a user might want to perform sixteen acquisitions of 1,100,000,000 Samples each using a 12-bit CompuScope. Because each 12-bit Sample occupies 2 data bytes, GageFastAcquire2Disk will only allow storage of 128 MegaSamples per SIG file. In this case, the 1.1 Billion Samples of data from a single acquisition will be divided across nine SIG files. The first eight files will contain 128 MegaSamples, and the ninth file will contain the remaining $1,100,000,000 - 8 \times 128 \times 1024 \times 1024 = 26,258,176$ Samples.

These nine SIG files will contain completely contiguous data so that, for example, the first Sample in the third SIG file was acquired directly after the last Sample in the second SIG file. The relationship amongst the nine contiguous SIG files is indicated by the Start Address contained within the SIG files. As is often the case, let us assume that the data requested for storage by the user began at the Trigger Address, which is always presented by the software as Trigger Address 0. In this case, the first SIG file will contain a Start Address of 0, which indicates that the first Sample within the file is the Trigger Point. The second SIG file, however, will contain a Start Address of $-128 \times 1024 \times 1024 = -134,217,728$, which correctly indicates that the Start Address occurred 134,217,728 before the first Sample in the second file. The Start Address will decrease by 134,217,728 Samples for each successive SIG file in the nine file set.

Since the user requested sixteen acquisitions of 1.1 Billion Samples each, each of the 16 acquisitions will create its own set of nine contiguous files. While each set of nine files represents a continuous acquisition within no breaks between files, between each file set there will be a time-gap as acquisition ceases while data are downloaded from the CompuScope hardware and stored to files.

In Multiple Record Mode, GageFastAcquire2Disk stores files in a similar fashion as in Single Record Mode. GageFastAcquire2Disk always tries to store all records from a single Multiple Record acquisition within a single SIG file. However, GageFastAcquire2Disk never mixes records from different Multiple Record acquisitions within the same SIG file. Division of data from a single Multiple Record is avoided, if possible. If the total Multiple Record data volume exceeds 256 MB, GageFastAcquire2Disk will distribute records across multiple SIG files, as in the first case above, while not splitting records if possible. The rare extreme Multiple Record case occurs when the user acquires very long Multiple Records. In this case, GageFastAcquire2Disk is forced to split the records across multiple files. For example, if a user chooses to acquire ten Multiple Records of 600,000,000 Samples each from an 8-bit CompuScope, GageFastAcquire2Disk will create ten contiguous SIG files sets – one set for each record. Each set will consist of three SIG Files, the first two of size 256 MB and the last one containing the roughly 100 MB of remaining left-over data.

The GageScope SIG files created by Gage FastAcquire2Disk contain an informative 512 Byte header, followed by the CompuScope waveform data in binary. Gage provides the `CsConvertFromSigHeader()` method, which decodes the header (from a buffer into which a SIG file header has transferred) and fills a structure with relevant header variables. These variables may then be used to select and interpret subsequent SIG file waveform data. Users can also use the `CsConvertToSigHeader()` method to create their own SIG file headers. Corresponding methods exist in the other SDKs.

Key elements that retard GageFastAcquire2Disk's maximum achievable repetitive acquisition rate are the time required to download data from the CompuScope hardware to PC RAM and the time to transfer data from PC RAM to the target hard drive. In almost all cases, it is the latter element that is the primary bottleneck to data throughput. Maximizing the storage rate to drive requires that certain optimization steps.

First of all, users should be wary of high storage rate claims by hard drive manufacturers. Usually, these claims refer to the data storage rate that may be maintained by the drive during burst storage. Most drive devices are equipped with a limited amount of cache memory. Usually, manufacturers' throughput claims refer only to burst storage volumes that do not exceed this cache size, which is typically a few MB. While very high storage rates may be realized into this disk cache memory, as soon as the cache is full, the rate drops dramatically, since the cache contents must be more slowly transferred to the actual drive media. What the Gage user really requires is this slower *sustained* storage rate, which is usually not provided by the hard drive manufacturer.

For best data throughput, the users should dedicate a specific drive as a storage target and should not store other files on this drive. Most importantly, the user should make sure not to store GageFastAcquire2Disk data on the same drive upon which the Windows Operating System or Page File is installed. The constant read/write housekeeping operations that Windows performs in normal course will seriously compromise the storage rate of Gage FastAcquire2Disk.

For best Disk performance, the user should obtain hard disk device drivers from the hard drive manufacturer, rather than using the generic drivers included in Windows, which typically will not provide the fastest throughput for your manufacturer's drive. Finally, disk throughput may be improved dramatically by configuring multiple hard drives in a RAID disk array. A RAID disk array may realize sustained data throughput rates of 100 MB/s and more.

GageEvents, GageCallback

These sample programs illustrate advanced synchronization techniques for multi-threaded applications. These techniques are essential to creating a complex application for real time data analysis or for operating multiple inter-related instruments. This is because these techniques allow a multi-threaded application to perform other tasks while CompuScope hardware is busy acquiring or transferring data without the usual need to poll its status. Without the need for polling, data acquisition and transfer do not tax the CPU, leaving it free to perform other operations. Nevertheless, these techniques add significant complexity to the overall application design, making it prone to errors such as thread deadlock. Consequently, usage of these techniques should not be considered unless they are truly required.

GageEvents is a multi-threaded sample program that illustrates usage of the notification events that can be assigned to specific operation of the CompuScope, allowing synchronization between different threads of execution. GageEvents uses the "end-of-busy" and "end-of-transfer" event notifications to trigger appropriate operations. In parallel, GageEvents processes older waveform data to determine the minimum and maximum points within the waveform. The handling of events that is illustrated within GageEvents is the recommended method of synchronization in a multi-threaded C application.

Some environments, such as Visual Basic and LabWindows/CVI, do not allow the programmer to create multi-threaded programs directly. These environments, however, do provide a functionality called "Callbacks", which allows the programmer to associate a callback function with notification of an event. The thread associated with the callback is then launched within the CompuScope driver. Use of callbacks has limitations. For instance, only one callback function may be executed at a time. Synchronization using callbacks is illustrated by the GageCallback sample program.

Technical Support

We offer technical support for all our Software Development Kits.

In order to serve you better, we have created a web-based technical support system that is available to you 24 hours a day.

By utilizing the internet to the fullest, we are able to provide you better than ever technical support without increasing our costs, thereby allowing us to provide you the best possible product at the lowest possible price.

To obtain technical support, simply visit:

<http://www.gage-applied.com/support/support-form.php>

Please complete this form and submit it. Our form processing system will intelligently route your request to the Technical Support Specialist (TSS) most familiar with the intricacies of your product. This TSS will be in contact with you within 24 hours of form submittal.

In the odd case that you have problems submitting the form on our web site, please e-mail us at

tech-support@gage-applied.com

As opposed to automatic routing of technical support requests originating from the GaGe web site, support requests received via e-mail or telephone calls are routed manually by our staff. Providing you with high quality support may take an average of 2 to 3 days if you do not use the web-based technical support system.

**Please note that Technical Support Requests received
via e-mail or by telephone will take an average of 2 to 3 days to process.
It is faster to use the web site!**

When calling for support we ask that you have the following information available:

1. Version and type of your CompuScope SDK and drivers.
(The version numbers are indicated in the About CD screen of the CompuScope CD. Version numbers can also be obtained by looking in the appropriate README.TXT files)
2. Type, version and memory depth of your CompuScope card.
3. Type and version of your operating system.
4. Type and speed of your computer and bus.
5. If possible, the file saved from the Information tab of the CompuScope Manager utility.
6. Any extra hardware peripherals (example: RAID Controller, DSP board, etc.)
7. Were you able to reproduce the problem with standalone Gage Software (e.g. GageScope, CsTest)?